

Assignment-1 [Total Marks - 25]

M1: MLOps Foundations

Objective: Understand the basics of MLOps and implement a simple CI/CD pipeline.

Tasks:

1. Set Up a CI/CD Pipeline:

- Use a CI/CD tool like GitHub Actions or GitLab CI to set up a pipeline for a sample machine learning project.
- Include stages for linting, testing, and deploying a simple machine learning model.

2. Version Control:

- Implement version control for your project using Git.
- Demonstrate branching, merging, and pull requests.

Deliverables:

- A report detailing the CI/CD pipeline stages.
- Screenshots or logs showing successful runs of the pipeline.
- A Git repository link with branches and merge history.

M2: Process and Tooling

Objective: Gain hands-on experience with popular MLOps tools and understand the processes they support.

Tasks:

1. Experiment Tracking:

- Use MLflow to track experiments for a machine learning project.
- Record metrics, parameters, and results of at least three different model training runs.

2. Data Versioning:

- Use DVC (Data Version Control) to version control a dataset used in your project.

- Show how to revert to a previous version of the dataset.

Deliverables:

- MLflow experiment logs with different runs and their results.
- A DVC repository showing different versions of the dataset.

M3: Model Experimentation and Packaging

Objective: Train a machine learning model, perform hyperparameter tuning, and package the model for deployment.

Tasks:

1. Hyperparameter Tuning:

- Use a library like Optuna or Scikit-learn's GridSearchCV to perform hyperparameter tuning on a chosen model.
- Document the tuning process and the best parameters found.

2. Model Packaging:

- Package the best-performing model using tools like Docker and Flask.
- Create a Dockerfile and a simple Flask application to serve the model.

Deliverables:

- A report on hyperparameter tuning results.
- A Dockerfile and Flask application code.
- Screenshots of the model running in a Docker container.

M4: Model Deployment & Orchestration (Optional)

Objective: Deploy a machine learning model and orchestrate its operations using Kubernetes.

Tasks:

1. Model Deployment:

- Deploy the Dockerized model from M3 to a cloud platform like AWS, Azure, or GCP.
- Use a platform service like AWS ECS, Azure AKS, or Google Kubernetes Engine (GKE).

2. Orchestration:

- Set up a Kubernetes cluster.
- Deploy the model using Kubernetes and create a Helm chart for managing deployments.

Deliverables:

- A link to the deployed model endpoint.
- Kubernetes configuration files and Helm chart.
- A report detailing the deployment and orchestration process.

M5: Final Deliverables

- A zip file containing:
 - Code
 - Data
 - Model
- A one-page summary that includes:
 - Description of the work completed
 - Justification for the choices made
- A screen recording (maximum 5 minutes) that:
 - Explains the work done
 - Shows the results